

# Enabling Partial Pivoting in Task Flow LU Factorization

Master Defense

Omar ZENATI

Supervisors: G. Bosilca, P. Ramet

September 13, 2012



- ▶ Three months at *Innovative Computer Laboratory*



- ▶ Three months at *Inria Bordeaux Sud-Ouest*



# Summary

Context and Motivations

LU Decomposition Algorithms

LU Decomposition over Task Flow Model

Performance

# Summary

Context and Motivations

LU Decomposition Algorithms

LU Decomposition over Task Flow Model

Performance

# Architecture Trends

- ▶ Computing platforms are more and more complex
- ▶ We classify them into four categories:
  - ▶ Shared memory architectures
  - ▶ Distributed memory architectures
  - ▶ Hierarchical architectures
  - ▶ Heterogeneous architectures

# Programming Paradigms

Historically, we use:

- ▶ Message passing (MPI) for distributed memory architectures
- ▶ Threads library (OMP) for shared memory architectures
- ▶ Accelerator library (Cuda, OpenCL) for accelerator of heterogeneous architectures

⇒ Several programs for the same algorithm

⇒ Weak portability

⇒ Weak scalability

Solution:

Other paradigm of programming which separate algorithm from architectures used: **Task Flow Model**

# Programming Paradigms

Historically, we use:

- ▶ Message passing (MPI) for distributed memory architectures
  - ▶ Threads library (OMP) for shared memory architectures
  - ▶ Accelerator library (Cuda, OpenCL) for accelerator of heterogeneous architectures
- ⇒ Several programs for the same algorithm
- ⇒ Weak portability
- ⇒ Weak scalability

## Solution:

Other paradigm of programming which separate algorithm from architectures used: **Task Flow Model**

# Programming Paradigms

Historically, we use:

- ▶ Message passing (MPI) for distributed memory architectures
  - ▶ Threads library (OMP) for shared memory architectures
  - ▶ Accelerator library (Cuda, OpenCL) for accelerator of heterogeneous architectures
- ⇒ Several programs for the same algorithm
- ⇒ Weak portability
- ⇒ Weak scalability

## Solution:


Other paradigm of programming which separate algorithm from architectures used: **Task Flow Model**



# Task Flow Model

Programs can be represented by a Direct Acyclic Graph (DAG) where:

- ▶ Vertices are tasks
- ▶ Edges are data dependencies between tasks



Execution of tasks and data movement between them assured by the **Runtime**

⇒ One single implementation of the DAG in a specific language

# Runtimes

## Advantages :

- ▶ Abstraction of architectures
- ▶ Portability of performance
- ▶ Good reactivity for load imbalance

## Question:

How to cope with parallel program consisting fine grain tasks?

## Challenge:

- ▶ DAGuE for large hierarchical architectures
- ▶ StarPU for heterogeneous architectures

# Summary

Context and Motivations

LU Decomposition Algorithms

LU Decomposition over Task Flow Model

Performance

# Why LU decomposition

In order to solve square systems of linear equations:

$$Ax = b$$

Use LU decomposition:

$$A = LU$$

Where  $L$  is a lower triangular matrix with the identity diagonal and  $U$  an upper triangular matrix.

Then solve:

$$Ly = b \text{ then } Ux = y$$

# LU algorithm

Let be  $A$  a  $n * n$  square matrix  
For  $k$  from 1 to  $n$


For  $i$  from  $k + 1$  to  $n$

$$a_{i,k} = a_{i,k} / a_{k,k}$$

For  $i$  from  $k + 1$  to  $n$

For  $j$  from  $k + 1$  to  $n$

$$a_{i,j} = a_{i,j} - a_{i,k} * a_{k,j}$$



# LU decomposition problem

Let be  $A$  a  $n * n$  square matrix

For  $k$  from 1 to  $n$

For  $i$  from  $k + 1$  to  $n$

$$a_{i,k} = a_{i,k} / a_{k,k}$$

For  $i$  from  $k + 1$  to  $n$

For  $j$  from  $k + 1$  to  $n$

$$a_{i,j} = a_{i,j} - a_{i,k} * a_{k,j}$$

- ▶  $a_{k,k}$  may be equal or close to zero
- ▶ Numerical value may be deteriorated due to fixed precision used by computers

⇒ LU decomposition is not stable

# LU decomposition problem

Let be  $A$  a  $n * n$  square matrix

For  $k$  from 1 to  $n$

For  $i$  from  $k + 1$  to  $n$

$$a_{i,k} = a_{i,k} / a_{k,k}$$

For  $i$  from  $k + 1$  to  $n$

For  $j$  from  $k + 1$  to  $n$

$$a_{i,j} = a_{i,j} - a_{i,k} * a_{k,j}$$

- ▶  $a_{k,k}$  may be equal or close to zero
- ▶ Numerical value may be deteriorated due to fixed precision used by computers

⇒ LU decomposition is not stable

# LU decomposition problem

Let be  $A$  a  $n * n$  square matrix

For  $k$  from 1 to  $n$

Search for pivot then swap

For  $i$  from  $k + 1$  to  $n$

$$a_{i,k} = a_{i,k} / a_{k,k}$$

For  $i$  from  $k + 1$  to  $n$

For  $j$  from  $k + 1$  to  $n$

$$a_{i,j} = a_{i,j} - a_{i,k} * a_{k,j}$$

- ▶  $a_{k,k}$  may be equal or close to zero
- ▶ Numerical value may be deteriorated due to fixed precision used by computers

⇒ LU decomposition is not stable

Solution is pivoting



# Partial Pivoting Algorithm

The partial pivoting consist to look for the element with the maximal absolute value on the  $k^{th}$  column from  $a_{k,k}$ , then swap its row with the row consisting  $a_{k,k}$ .

The factorization amount to  $PA = LU$


The partial pivoting is:

- ▶ Practically stable and accurate
- ▶ Commonly used in the scientific community
- ▶ Used in the LINPACK benchmark to rank the TOP 500 super-computers

# LU implementation ( $PA = LU$ )

Software/Algorithms follow hardware evolution in time


- ▶ 70's - LinPACK, vector operations
- ▶ 80's - LAPACK, block, cache-friendly
- ▶ 90's - ScaLAPACK, distributed memory



# LU implementation ( $PA = LU$ )

Software/Algorithms follow hardware evolution in time


- ▶ 70's - LinPACK, vector operations
- ▶ 80's - LAPACK, block, cache-friendly
- ▶ 90's - ScaLAPACK, distributed memory



# LU implementation ( $PA = LU$ )

Software/Algorithms follow hardware evolution in time

- ▶ 70's - LinPACK, vector operations
- ▶ 80's - LAPACK, block, cache-friendly
- ▶ 90's - ScaLAPACK, distributed memory



# Summary

Context and Motivations


LU Decomposition Algorithms

LU Decomposition over Task Flow Model

Performance


## Task flow LU ( $A = LU$ )

In order to cope with the task flow model, linear algebra algorithms are expressed in terms of tasks operating on fine grain squares sub-matrices, also called tiles.



## Task flow LU ( $A = LU$ )


In order to cope with the task flow model, linear algebra algorithms are expressed in terms of tasks operating on fine grain squares sub-matrices, also called tiles.



Tiles are not contiguous in memory

## Task flow LU ( $A = LU$ )

In order to cope with the task flow model, linear algebra algorithms are expressed in terms of tasks operating on fine grain squares sub-matrices, also called tiles.




GETRF(0)

Task to factorize diagonal tiles: GETRF



## Task flow LU ( $A = LU$ )



In order to cope with the task flow model, linear algebra algorithms are expressed in terms of tasks operating on fine grain squares sub-matrices, also called tiles.



Task to update other panel tiles and block line tiles: TRSM\_L and TRSM\_U

# Task flow LU ( $A = LU$ )



In order to cope with the task flow model, linear algebra algorithms are expressed in terms of tasks operating on fine grain squares sub-matrices, also called tiles.



Task to update trailing sub-matrix: GEMM


# Task flow LU ( $A = LU$ )

In order to cope with the task flow model, linear algebra algorithms are expressed in terms of tasks operating on fine grain squares sub-matrices, also called tiles.



# Panel Factorization Problems


## Tasks of panel factorization



- Look for a pivot is distributed over several tiles
- Tasks are fine grained

# Panel Factorization Problems

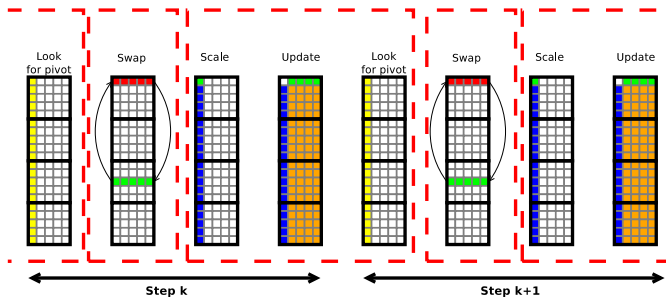
## Tasks of panel factorization



- ▶ Look for a pivot is distributed over several tiles
- ▶ Tasks are fine grained


# Panel Factorization Problems

Reduce fine grained tasks



# Panel Factorization Problems

## Natural task flow of panel factorization




⇒ Serialized task flow!

Optimizations:

Use *all\_reduce* operation (using Bruck's algorithm)

# Panel Factorization Problems

## Natural task flow of panel factorization



⇒ Serialized task flow!


Optimizations:

Use *all\_reduce* operation (using Bruck's algorithm)



# Panel Factorization Problems

## Natural task flow of panel factorization



⇒ Serialized task flow!

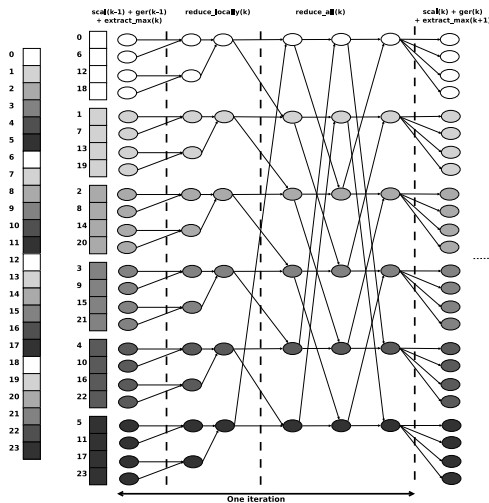
## Optimizations:

Use *all\_reduce* operation (using Bruck's algorithm)




# Panel Factorization Problems

Optimized task flow of panel factorization for hierarchical architectures



# Update Problems

## Update trailing sub-matrix




The upper tile exchange rows with other tile depending on pivots values.

## Problem

- ▶ Dynamic decision for a static DAG
- ▶ Pivots implies swaps in a specific order
- ▶ Serialized communications

# Update Problems

## Update trailing sub-matrix




The upper tile exchange rows with other tile depending on pivots values.

## Problem

- Dynamic decision for a static DAG
- Pivots implies swaps in a specific order
- Serialized communications

# Update Problems

## Update trailing sub-matrix



The upper tile exchange rows with other tile depending on pivots values.

## Problem → Solutions

- Dynamic decision for a static DAG → Generate tasks for all possible communications?
- Pivots implies swaps in a specific order → Use permutation instead of pivots
- Serialized communications → Separate Swap from/into upper tile

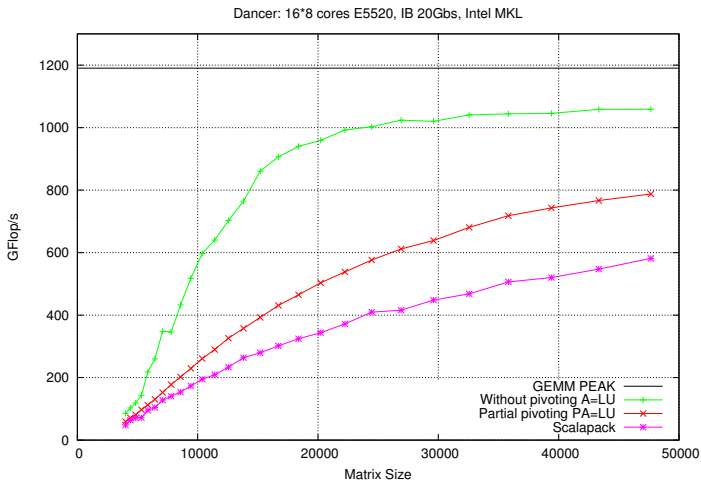
# Summary

Context and Motivations

LU Decomposition Algorithms

LU Decomposition over Task Flow Model

Performance





# Conclusion and future work

## Software contribution:

- ▶ Implemented and validated  $A = LU$  on DAGuE and StarPU
- ▶ Implemented and validated  $PA = LU$  on DAGuE and StarPU

## Conclusion:

- ▶ Exhibited the feasibility of partial pivoting on task flow model
- ▶ Obtained encouraging performance with partial pivoting over DAGuE

## Future work:


- ▶ Integrate GPU pivoting operations to DAGuE and StarPU
- ▶ Try other strategies of panel factorizations on several methods
- ▶ Build a new benchmark based on the DAGuE partial pivoting

# Thank you !

# ANNEXE

## Using permutations instead of pivots


### Using pivots




# ANNEXE

## Using permutations instead of pivots

### Using pivots




### Using permutations




# Update Problems

## Natural task flow of one swap in update




# Update Problems

Optimized task flow of swaps of update for distributed architectures




# Update Problems

Optimized task flow of swaps of update for distributed architectures



# Update Problems


Optimized task flow of swaps of update for distributed architectures





# Update Problems

Optimized task flow of swaps of update for distributed architectures



# Update Problems

## Optimized task flow of swaps of update for hierarchical architectures

